



Modality geometry across omni-style embedding models — UMAP on MSR-VTT parallel (image, video, audio, text) tuples


jina-v5-omni-small-retrieval
(1.57 B)




jina-v5-omni-nano-retrieval
(0.95 B)




LCO-Omni-3B
(3.7 B)




LCO-Omni-7B
(8.93 B)



Omni-Embed-Nemotron-3B
(4.7 B)



LanguageBind
(0.43-0.53 B/modality)



• image • video • audio • text