



(a) Training Loss Gap vs. MLA



(b) Validation Loss (128K Context)



(c) Needle-in-a-Haystack Evaluation

